The Agent Register

Every agent that can act: counted, owned, scoped, and stoppable

Version 1.0 · October 2026 · Nitin Gupta

A team can now stand up an AI agent in an afternoon. Every agent that can read a system or take an action holds some authority, whether or not anyone granted it. An agent nobody inventoried has authority nobody approved, an owner nobody named, and a stop procedure nobody has tested.

The rest of this toolkit governs systems the enterprise knows about. The Agent Register makes sure that is all of them. It has three parts: a **Scope Card** completed for each agent before it goes live, the **Register** that holds every card in one place, and a quarterly **Reconciliation** that compares the register with what is actually running.

Three rules sit underneath it:

1. **Every agent needs three things before it goes live: a scope, an owner, and a stop button.** The scope says what it may touch and what it may never touch. The owner is one named person who answers for it. The stop button is the tested way to take its authority back.
2. **No entry, no credentials.** An agent that is not in the register does not receive a service identity, an API key, or a write permission. Identity and access teams enforce this rule, not the agent's builders.
3. **The register is only true if it is reconciled.** Compare it with identities, keys, and scheduled jobs at least quarterly. A register nobody reconciles describes the agents someone remembered.

What counts as an agent

Register any software that can take an action in an enterprise system on the basis of a model's output: write a record, send a message, approve a transaction, change a setting, schedule a job, or call another agent that does any of these. Two tiers keep the register usable.

Tier	What it covers	What the register needs
Acting	Can write to, send from, or change anything in an enterprise system, directly or through another agent or tool	A full Scope Card
Read-only	Reads enterprise data and answers a person, with no ability to change anything	Identity, owner, data sources, and tier. Upgrade to Acting before any write permission is granted

If in doubt, it is Acting. An assistant that can draft an email is Read-only. One that can send it is Acting.

Part One: The Scope Card

Complete before go-live. The card becomes the agent's register entry, and its first Authority Change Record refers to it.

Identity

Field	Answer
Agent ID	Unique, never reused
Name and purpose	One sentence: what decision or task it supports
Tier	Acting / Read-only
Non-human identity	The service identity or account it runs as, distinct from any person's
Built or bought	Built / Bought / Assembled. If bought, the vendor and product
Model and version	Pinned version, and who approves a change
Markets and business units	Where it runs

Scope: what it may touch

Field	Answer
Systems it may read	Each system, with the data classification
Systems it may write to	Each system, with the action types allowed in each
Limits	Numeric caps per action type: value, volume, records changed per hour
May never touch	An explicit list: systems, data, action types, and customers or counterparties that are out of bounds, even if access exists
Decisions it may prepare	Analysis, drafts, and recommendations a person decides on
Decisions it may make	Actions it takes without a person approving each one, if any. Each must match its authority state
Agents and tools it may call	Every delegation path. No path may grant more than this card allows

The “may never touch” line is the one teams skip and the one that matters most. Access drifts. An explicit exclusion list is how a reviewer notices when it has.

Owner: who answers for it

Role	Name	Holds
Accountable decision owner		The right to approve, expand, restrict, or revoke its authority
Business owner		The outcome and economics it affects
Technical owner		Operation, monitoring, and recovery
Backup for the decision owner		The same powers when the owner is unavailable

One name per row. A team or a committee is not an owner.

Stop button: how its authority is taken back

Field	Answer
How to stop it	The mechanism: a switch, a revoked credential, a disabled policy. Not a code deployment
Who may press it	Named people, including someone available outside business hours
Scope of the stop	All actions, or also by action type and market
Time to stop, as tested	Minutes from decision to the last state-changing action
Last tested	Date, and the record of the test
What happens to work in flight	Actions started but not finished: completed, held, or cancelled, and by whom
Fallback	Who does the work while it is stopped

An untested stop button is a belief, not a control. If you cannot take an agent’s authority back, you never really gave it. You lost it.

Authority at go-live

Action type	Authority state	Limits	Authority Change Record
	Observe / Recommend / Act with approval / Bounded autonomy / Expanded autonomy		

New agents start at Observe or Recommend. Anything higher at go-live needs a completed AI Authority Review and a Scorecard band that allows it.

Signatures

Role	Name	Date
Accountable decision owner		
Technical owner (confirms identity, limits, and stop button are configured as written)		
Identity and access (confirms credentials match this card)		

Part Two: The Register

One row per agent, in one place, for the whole enterprise. A CSV template with these columns is included, with one fictional example row to show the format.

Column	What it holds
Agent ID	From the Scope Card
Name	Short name
Purpose	One sentence
Tier	Acting / Read-only
Non-human identity	Service identity or account
Built or bought	And vendor, if bought
Model and version	Pinned version
Markets	Where it runs
Systems written to	From the Scope Card
Highest authority state	The highest state held for any action type
Accountable decision owner	One name
Business owner	One name
Technical owner	One name
Stop button last tested	Date
Latest Authority Change Record	Record number
Authority sunset	Date the current grant expires, if any
Scorecard band and date	From the most recent Readiness Scorecard
Next review	Date
Reconciliation status	Matched / Drifted / Unregistered / Retired, from the last sweep

Three columns earn their place more than the rest. **Stop button last tested** shows how many agents could be stopped today. **Authority sunset** shows how many grants are quietly lapsing without review. **Highest authority state**, read against **Scorecard band**, shows how many agents hold more authority than their evidence allows.

Part Three: The Reconciliation

At least quarterly, and after any reorganization, platform migration, or acquisition. Owned by the identity and access team with the AI governance lead.

Step 1: Pull what is actually running. From the identity provider, every non-human identity with write access to an enterprise system. From each AI and automation platform, every agent, API key, and scheduled job. From procurement, every product bought since the last sweep that includes agent features.

Step 2: Match it to the register. Each running agent lands in one of four outcomes.

Outcome	What it means	Required response	Deadline
Matched	Registered, and its real permissions match its Scope Card	None	
Drifted	Registered, but its real permissions exceed its Scope Card, or its owner has left	Restrict permissions to the card, or raise an Authority Change Record to change the card. Replace any departed owner	10 business days
Unregistered	Running, with credentials, and not in the register	Suspend its write permissions now. Complete a Scope Card, or retire it	Suspend at once. Card or retire within 10 business days
Retired	In the register, but no longer running or needed	Revoke every credential and close the entry with a dated note	10 business days

Step 3: Report four numbers to the AI governance forum.

1. Unregistered agents found this sweep, and how many were acting agents.
2. Share of acting agents whose stop button was tested in the last 90 days.
3. Share of acting agents holding more authority than their latest Scorecard band allows.
4. Authority grants that passed their sunset without renewal.

The first number is the one to watch over time. It should fall with every sweep. If it does not, the “no entry, no credentials” rule is not being enforced.

How to use it

- **Make the register the gate for credentials.** The single rule that does the most work is that identity and access will not issue an agent credential without a register entry. Everything else follows from it.
- **Start with the agents you have.** Run the first reconciliation before writing any new cards. Most enterprises find agents they did not know about, and those are the ones to card first.
- **Keep it boring.** A spreadsheet or an existing configuration management database is enough. The register fails when it is incomplete, not when it is simple.
- **Bring it to the scheduled authority review.** The register shows which agents are due for review. The AI Authority Review decides. The Authority Change Record records it.

Relationship to the Scorecard

The Register is how a team earns Scorecard criteria 7.4 (agent identity and inventory), 2.1 and 2.2 (data sources and least-privilege access), 5.1 (tested kill switch), and 5.5 (delegation controls). A Scorecard applies to one system. The Register applies to all of them.

Lineage

Asset inventories and configuration management databases are standard IT practice, and access recertification is a standard identity governance control. Banking supervisors have long expected firms to keep an inventory of their models, as in the US Federal Reserve and OCC guidance on model risk management (SR 11-7, 2011). What the Agent Register adds for AI is a required scope with an explicit “may never touch” list, the authority state and a tested stop button as required fields for every entry, the rule that no entry means no credentials, and a reconciliation designed to find agents that were never registered at all.

Part of the Enterprise AI Decision Toolkit, a set of open instruments for enterprise AI governance based on *The Enterprise AI Decision System: How AI Leaders Turn Capability into Authorized Enterprise Action* by Nitin Gupta. Latest version: github.com/ngupta1137/enterprise-ai-decision-toolkit. Licensed under MIT.